



Assignment 1 (29.04.2022)

Handin until: 06.05.2022, 09:00

1. [0 Points] Introduction

- (a) **Before we are able to grade you** and/or your team, you have to complete this initial task first. In your team's GitHub Classroom repository, there exists a file called `README.md`. Add both team member names, surname, matrikel number, etc. to the file. This task does not grant you any points, but is a **requirement** for your team to be graded in the first place. This file has to persist throughout the entire semester.
- (b) Each submission of each assignment in "Advanced SQL" must be put into its own folder in the **root** directory of your team's GitHub Classroom repository. The name of the folder has to be in the format `solution<number>`. For example: the folder of your submission for this assignment must be called `solution01`.
- (c) In general, the only accepted file format is plain text (*.txt, *.sql for SQL code). Other file formats are now allowed, unless stated otherwise. Your submitted code has to be consistently formatted, well-documented and has to run without errors using PostgreSQL 14. Older versions of PostgreSQL may also work, but for those you are on your own.
- (d) Lastly, the usual rules for plagiarism and other academic integrity apply.
- (e) For further questions about this lecture, assignment and other related topics, visit the forum at

<https://forum-db.informatik.uni-tuebingen.de/c/ss22-asql>.

2. [5 Points] So similar, yet so different

Create an instance for table `r` with schema `r(a int, b int)` such that the two queries Q1 and Q2 compute **different** results.

<pre> 1 -- Q1 2 SELECT r.a, COUNT(*) AS c 3 FROM r AS r 4 WHERE r.b <> 3 5 GROUP BY r.a; </pre>	<pre> 1 -- Q2 2 SELECT r.a, COUNT(*) AS c 3 FROM r AS r 4 GROUP BY r.a 5 HAVING EVERY(r.b <> 3); </pre>
---	--

Note: Query Q2 uses the aggregate function `EVERY`. Read about it in the PostgreSQL documentation¹.

3. [10 Points] Production Steps

Consider table `production` which tracks the progress of items (column `item`) currently in production. Production occurs in multiple steps (column `step`). If a step has been completed for an item, column `completion` indicates the time and date. If the step is still in progress, column `completion` holds `NULL` for that step. Table `production` is defined as follows:

```

1  CREATE TABLE production (
2  item          char(20) NOT NULL,
3  step          int      NOT NULL,
4  completion    timestamp,      -- NULL means incomplete
5  PRIMARY KEY (item, step));

```

¹<https://www.postgresql.org/docs/current/functions-aggregate.html>

For **example**, in this instance

product_name	step	completion_date
'TIE'	1	'1977/03/02 04:12'
'TIE'	2	'1977/12/29 05:55'
'AT-AT'	1	'1978/01/03 14:12'
'AT-AT'	2	NULL
'DSII'	1	NULL
'DSII'	2	'1979/05/26 20:05'
'DSII'	3	'1979/04/04 17:12'

step 2 of the production of AT-AT is not complete yet. Likewise, Step 1 of DS II is pending, while Steps 2 and 3 are complete.

Formulate a SQL query that lists the names of all **items** for which all production steps are complete. Avoid duplicate item names. For the instance above, we expect the following result:

complete
'TIE'

4. [15 Points] Matrix Multiplication

You are given two tables representing two matrices A and B of sizes $m \times n$ and $n \times p$, respectively. We create those tables with three columns, where **row** represents the row index and **col** represents the column index of the matrix. The column **val** represents the value of a matrix element.

```

1 CREATE TABLE A (
2   row int,
3   col int,
4   val int,
5   PRIMARY KEY(row, col));
1 CREATE TABLE B ( LIKE A );

```

(a) Formulate a SQL query, which performs matrix multiplication $A \cdot B$.

Example:

You are given $A = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}$ and $B = \begin{pmatrix} 1 & 2 & 1 \\ 2 & 1 & 2 \end{pmatrix}$, so $A \cdot B = \begin{pmatrix} 5 & 4 & 5 \\ 11 & 10 & 11 \end{pmatrix}$.

So, for these instances

```

1 INSERT INTO A (row,col,val)
2 VALUES (1,1,1), (1,2,2),
3         (2,1,3), (2,2,4);
1 INSERT INTO B (row,col,val)
2 VALUES (1,1,1), (1,2,2), (1,3,1),
3         (2,1,2), (2,2,1), (2,3,2);

```

we expect the following result from your query:

row	col	val
1	1	5
1	2	4
1	3	5
2	1	11
2	2	10
2	3	11

(b) In *sparse matrices*, all entries of value 0 are missing. See the examples of sparse matrices A and B below. Does your SQL query of (a) also apply to sparse matrices?

```

1 INSERT INTO A (row,col,val)
2 VALUES (1,1,1), (1,2,3),
3         (2,3,7);
1 INSERT INTO B (row,col,val)
2 VALUES (1,1,4),           (1,3,8 ),
3         (2,1,1), (2,2,1), (2,3,10),
4         (3,1,3), (3,2,6);

```